

Homework 1

Kevin Yang - 50244152

September 30, 2021

Exercise 1

Let $X = (x_1, x_2, \dots, x_n)$

$$x_i \sim \text{Pois}(\lambda) \quad \lambda \sim \Gamma(\alpha, \beta)$$

We have that:

$$\begin{aligned} P(\lambda \mid X) &= \frac{P(X \mid \lambda)P(\lambda)}{P(X)} \\ &\propto P(X \mid \lambda)P(\lambda) \\ &\propto \lambda^{n\bar{x}} e^{-n\lambda} \lambda^{\alpha-1} e^{-\beta\lambda} \\ &= \lambda^{n\bar{x}+\alpha-1} e^{-(n+\beta)\lambda} \end{aligned}$$

Therefore, $P(\lambda \mid X) \sim \Gamma(n\bar{x} + \alpha, \beta + n)$ and the Gamma distribution is conjugate to the Poisson distribution.

Exercise 2

Let $s = (s_1, s_2, \dots, s_n)$, $s' = (s'_1, s'_2, \dots, s'_n)$ and $s_{-i} = (s_1, \dots, s_{i-1}, s_{i+1}, \dots, s_n)$

We will show that Gibbs sampling satisfies the detailed balance equation: $\pi(s)P(s, s') = \pi(s')P(s', s)$

w.l.g consider an update for s_1

Case 1: $s_{-1} \neq s'_{-1}$

$$\pi(s)P(s, s') = \pi(s')P(s', s) = 0$$

Case 2: $s_{-1} = s'_{-1}$

$$\begin{aligned}\pi(s)P(s, s') &= \pi(s)P(s'_1 \mid s_{-1}) \\ &= \pi(s) \frac{\pi(s')}{\sum_z \pi(z, s'_{-1})} \\ &= \pi(s') \frac{\pi(s)}{\sum_z \pi(z, s_{-1})} \\ &= \pi(s')P(s_1 \mid s_{-1}) \\ &= \pi(s')P(s', s)\end{aligned}$$

Consider the move from s to s' . The acceptance probability for MH will be

$$\frac{\pi(s')P(s', s)}{\pi(s)P(s, s')} = 1$$

Exercise 3

a)

```
##1. enumeration and conditioning:

## condition and marginalize:
## compute joint:
p = np.zeros((2, 2, 2, 2)) # c,s,r,w
for c in range(2):
    for s in range(2):
        for r in range(2):
            for w in range(2):
                p[c, s, r, w] = p_C(c) * p_S_given_C(s, c) * p_R_given_C(r, c) * p_W_given_S_R(w, s, r)

p_C_given_W = np.zeros(2)
for c in range(2):
    for s in range(2):
        for r in range(2):
            p_C_given_W[c] += p[c, s, r, 1]

p_C_given_W /= np.sum(p_C_given_W)

print('There is a {:.2f}% chance it is cloudy given the grass is wet'.format(p_C_given_W[1] * 100))
```

b)

```
##2. ancestral sampling and rejection:
# https://www.cs.ubc.ca/~fwood/CS532W-539W/lectures/mcmc.pdf

num_samples = 10000
samples = np.zeros(num_samples)
rejections = 0
i = 0
while i < num_samples:
    c = np.argmax(np.random.multinomial(1, [p_C(0), p_C(1)]))
    s = np.argmax(np.random.multinomial(1, [p_S_given_C(0, c), p_S_given_C(1, c)]))
    r = np.argmax(np.random.multinomial(1, [p_R_given_C(0, c), p_R_given_C(1, c)]))
    w = np.argmax(np.random.multinomial(1, [p_W_given_S_R(0, s, r), p_W_given_S_R(1, s, r)]))
    if w != 1:
        rejections += 1
        continue
    else:
        samples[i] = c
        i += 1

print('The chance of it being cloudy given the grass is wet is {:.2f}%'.format(samples.mean() * 100))
print(' {:.2f}% of the total samples were rejected'.format(100 * rejections / (samples.shape[0] + rejections)))
```

c)

```

##gibbs sampling
num_samples = 10000
samples = np.zeros(num_samples)
state = np.zeros(4, dtype='int')
# c,s,r,w, set w = True

c, s, r, w = 0, 1, 2, 3
i = 0
state[w] = 1
while i < num_samples:
    state[c] = np.argmax(np.random.multinomial(1, p_C_given_S_R[:, state[s], state[r]]))
    state[s] = np.argmax(np.random.multinomial(1, p_S_given_C_R_W[state[c], :, state[r], state[w]]))
    state[r] = np.argmax(np.random.multinomial(1, p_R_given_C_S_W[state[c], state[s], :, state[w]]))

    samples[i] = state[c]
    i += 1

print('The chance of it being cloudy given the grass is wet is {:.2f}%'.format(samples.mean() * 100))

```

Results:

```

There is a 57.58% chance it is cloudy given the grass is wet
The chance of it being cloudy given the grass is wet is 58.05%
34.73% of the total samples were rejected
The chance of it being cloudy given the grass is wet is 58.91%

Process finished with exit code 0

```

Exercise 4

MH within Gibbs on blocks \mathbf{w} and $\hat{\mathbf{t}}$:

\mathbf{w} block

Let $q(\mathbf{w}, \mathbf{w}')$ be the proposal distribution. We also have that,

$$\begin{aligned}\pi(\mathbf{w}) &\propto p(\mathbf{t}, \mathbf{x}, \sigma^2, \mathbf{w}, \alpha) \\ &= \prod_{n=1}^N \frac{1}{\sigma \sqrt{2\pi}} e^{-\frac{1}{2}(\frac{t_n - \mathbf{w}^T \mathbf{x}_n}{\sigma})^2} \prod_{d=1}^D \frac{1}{\sqrt{2\pi\alpha}} e^{-\frac{1}{2}(\frac{\mathbf{w}_d}{\sqrt{\alpha}})^2}\end{aligned}$$

So the update probability is,

$$\begin{aligned}r &= \frac{\pi(\mathbf{w}')q(\mathbf{w}', \mathbf{w})}{\pi(\mathbf{w})q(\mathbf{w}, \mathbf{w}')} \\ &= \frac{\prod_{n=1}^N \frac{1}{\sigma \sqrt{2\pi}} e^{-\frac{1}{2}(\frac{t_n - \mathbf{w}'^T \mathbf{x}_n}{\sigma})^2} \prod_{d=1}^D \frac{1}{\sqrt{2\pi\alpha}} e^{-\frac{1}{2}(\frac{\mathbf{w}'_d}{\sqrt{\alpha}})^2} q(\mathbf{w}', \mathbf{w})}{\prod_{n=1}^N \frac{1}{\sigma \sqrt{2\pi}} e^{-\frac{1}{2}(\frac{t_n - \mathbf{w}^T \mathbf{x}_n}{\sigma})^2} \prod_{d=1}^D \frac{1}{\sqrt{2\pi\alpha}} e^{-\frac{1}{2}(\frac{\mathbf{w}_d}{\sqrt{\alpha}})^2} q(\mathbf{w}, \mathbf{w}')} \\ &\propto \frac{\prod_{n=1}^N e^{-\frac{1}{2}(\frac{t_n - \mathbf{w}'^T \mathbf{x}_n}{\sigma})^2} e^{-\frac{1}{2}(\mathbf{w}'^T (\alpha I)^{-1} \mathbf{w}')} q(\mathbf{w}', \mathbf{w})}{\prod_{n=1}^N e^{-\frac{1}{2}(\frac{t_n - \mathbf{w}^T \mathbf{x}_n}{\sigma})^2} e^{-\frac{1}{2}(\mathbf{w}^T (\alpha I)^{-1} \mathbf{w})} q(\mathbf{w}, \mathbf{w}')} \\ &= \frac{e^{-\frac{1}{2}(\mathbf{t} - \mathbf{x} \mathbf{w}')^T (\sigma^2 I)^{-1} (\mathbf{t} - \mathbf{x} \mathbf{w}') - \mathbf{w}'^T (\alpha I)^{-1} \mathbf{w}'} q(\mathbf{w}' | \mathbf{w}')}{e^{-\frac{1}{2}(\mathbf{t} - \mathbf{x} \mathbf{w})^T (\sigma^2 I)^{-1} (\mathbf{t} - \mathbf{x} \mathbf{w}) - \mathbf{w}^T (\alpha I)^{-1} \mathbf{w}} q(\mathbf{w} | \mathbf{w})}\end{aligned}$$

$\hat{\mathbf{t}}$ block

Let $q(\mathbf{w}, \mathbf{w}')$ be the proposal distribution.

$$\pi(\hat{\mathbf{t}}) \propto N(\mathbf{w}^T \hat{\mathbf{x}}, \sigma^2)$$

The update probability is,

$$\begin{aligned}r &= \frac{\pi(\hat{\mathbf{t}}')q(\hat{\mathbf{t}}', \hat{\mathbf{t}})}{\pi(\hat{\mathbf{t}})q(\hat{\mathbf{t}}, \hat{\mathbf{t}}')} \\ &\propto \frac{e^{-\frac{1}{2}(\frac{\hat{\mathbf{t}}' - \mathbf{w}^T \hat{\mathbf{x}}}{\sigma})^2} q(\hat{\mathbf{t}}', \hat{\mathbf{t}})}{e^{-\frac{1}{2}(\frac{\hat{\mathbf{t}} - \mathbf{w}^T \hat{\mathbf{x}}}{\sigma})^2} q(\hat{\mathbf{t}}, \hat{\mathbf{t}}')}\end{aligned}$$

Pure Gibbs on blocks \mathbf{w} and $\hat{\mathbf{t}}$:

\mathbf{w} block

$$\begin{aligned}
\log p(\mathbf{w} \mid \mathbf{t}, \mathbf{x}, \sigma^2, \alpha) &\propto \log p(\mathbf{w}, \mathbf{t}, \mathbf{x}, \sigma^2, \alpha) \\
&= \log \prod_{n=1}^N p(t_n \mid \mathbf{w}, \mathbf{x}_n, \sigma^2) p(\mathbf{w} \mid 0, \alpha \mathbf{I}) \\
&\propto -\frac{\sum_{n=1}^N (t_n - \mathbf{w}^T \mathbf{x}_n)^2}{2\sigma^2} - \frac{1}{2} \mathbf{w}^T (\alpha \mathbf{I})^{-1} \mathbf{w} \\
&\propto -\frac{1}{2\sigma^2} (\|\mathbf{t} - \mathbf{x}^T \mathbf{w}\|^2) - \frac{1}{2} \mathbf{w}^T (\alpha \mathbf{I})^{-1} \mathbf{w} \\
&= -\frac{1}{2\sigma^2} (\mathbf{t}^T \mathbf{t} - 2\mathbf{t}^T \mathbf{x}^T \mathbf{w} + \mathbf{w}^T \mathbf{x} \mathbf{x}^T \mathbf{w}) - \frac{1}{2} \mathbf{w}^T (\alpha \mathbf{I})^{-1} \mathbf{w} \\
&= -\frac{1}{2} (\mathbf{w} - \boldsymbol{\mu}) \boldsymbol{\Sigma}^{-1} (\mathbf{w} - \boldsymbol{\mu}) \sim N(\boldsymbol{\mu}, \boldsymbol{\Sigma}) \\
\boldsymbol{\mu} &= \sigma^{-2} \boldsymbol{\Sigma} \mathbf{x}^T \mathbf{t} \qquad \boldsymbol{\Sigma}^{-1} = \sigma^{-2} \mathbf{x} \mathbf{x}^T + (\alpha \mathbf{I})^{-1}
\end{aligned}$$

$\hat{\mathbf{t}}$ block

$$p(\hat{\mathbf{t}} \mid \mathbf{t}, \mathbf{x}, \sigma^2, \mathbf{w}, \alpha) \propto e^{-\frac{1}{2}(\frac{\hat{\mathbf{t}} - \mathbf{w}^T \hat{\mathbf{x}}}{\sigma})^2} \sim \text{Normal}(\mathbf{w}^T \hat{\mathbf{x}}, \sigma^2)$$

Posterior Predictive:

$$\begin{aligned}
p(\hat{\mathbf{t}} \mid \hat{\mathbf{x}}, \mathbf{t}) &= \int p(\hat{\mathbf{t}} \mid \hat{\mathbf{x}}, \mathbf{w}) p(\mathbf{w} \mid \mathbf{t}) d\mathbf{w} \\
&= \int N(\mathbf{w}^T \hat{\mathbf{x}}, \sigma^2) N(\mathbf{w} \mid \boldsymbol{\mu}, \boldsymbol{\Sigma}) d\mathbf{w}
\end{aligned}$$

Bye the linear combination rules for Gaussian random variables, we have that,

$$p(\hat{\mathbf{t}} \mid \hat{\mathbf{x}}, \mathbf{t}) \sim N(\boldsymbol{\mu}', \sigma'^2)$$

$$\begin{aligned}
\boldsymbol{\mu}' &= \boldsymbol{\mu}^T \hat{\mathbf{x}} \\
\sigma'^2 &= \hat{\mathbf{x}}^T \boldsymbol{\Sigma} \hat{\mathbf{x}} + \sigma^2
\end{aligned}$$

Exercise 5

M = number of documents, K = number of topics, V = vocabulary size

N_{wi} = number of times word w is assigned to topic i

N_i = number of words assigned to topic i

N_{ji} = number of words in document j assigned to topic i

$x_{lj} = l^{th}$ word in document j (observed)

z_{lj} = topic assignment for the l^{th} word in document j

$\hat{\phi}_i$ = distribution of words for topic i

$\hat{\theta}_j$ = distribution of topics for document j

Joint log likelihood:

$$\begin{aligned} \log p(z, w \mid \alpha, \beta) \propto & \sum_{j=1}^M \left(\sum_{i=1}^K \log \Gamma(N_{ji} + \alpha_i) \right) - \log \Gamma\left(\sum_{i=1}^K N_{ji} + \alpha_i\right) \\ & + \sum_{i=1}^K \left(\sum_{w=1}^V \log \Gamma(N_{wi} + \beta_w) \right) - \log \Gamma\left(\sum_{w=1}^V N_{wi} + \beta_w\right) \end{aligned}$$

Conditional:

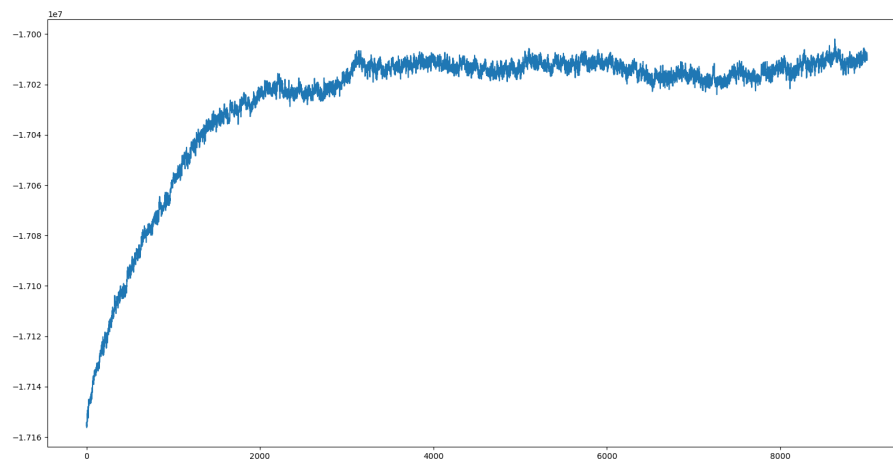
$$p(z_{lj} = k \mid z^{-lj}, x, \alpha, \beta) = \frac{1}{Z} a_{ji} b_{wi}$$

$$a_{ji} = N_{ji}^{-lj} + \alpha \qquad b_{wi} = \frac{N_{wi}^{-lj} + \beta}{N_i^{-lj} + V\beta} \qquad Z = \sum_i^K a_{ji} b_{wi}$$

Estimates:

$$\begin{aligned} \hat{\phi}_{wi} &= \frac{N_{wi} + \beta}{N_i + V \cdot \beta} \\ \hat{\theta}_{ji} &= \frac{N_{ji} + \alpha}{N_j + K \cdot \alpha} \end{aligned}$$

The joint log likelihood plot,



The 10 most probable words per topic (rows correspond topics)

```
['robot'] ['trajectory'] ['head'] ['eye'] ['position'] ['control'] ['motor'] ['system'] ['model'] ['controller']
['neural'] ['weights'] ['networks'] ['training'] ['layer'] ['output'] ['input'] ['units'] ['hidden'] ['network']
['temporal'] ['neurons'] ['neuron'] ['rate'] ['frequency'] ['time'] ['spike'] ['information'] ['firing'] ['signal']
['feature'] ['vision'] ['figure'] ['objects'] ['motion'] ['features'] ['object'] ['images'] ['visual'] ['image']
['actions'] ['reinforcement'] ['time'] ['state'] ['function'] ['action'] ['states'] ['policy'] ['optimal'] ['learning']
['case'] ['networks'] ['class'] ['set'] ['threshold'] ['number'] ['bound'] ['theorem'] ['functions'] ['function']
['matrix'] ['approximation'] ['method'] ['function'] ['functions'] ['optimal'] ['linear'] ['error'] ['vector'] ['noise']
['speaker'] ['state'] ['training'] ['context'] ['hmm'] ['time'] ['speech'] ['word'] ['recognition'] ['system']
['classifier'] ['class'] ['neural'] ['performance'] ['error'] ['data'] ['test'] ['training'] ['set'] ['classification']
['model'] ['recurrent'] ['neuron'] ['system'] ['neurons'] ['state'] ['networks'] ['neural'] ['time'] ['network']
['values'] ['work'] ['figure'] ['data'] ['set'] ['problem'] ['point'] ['space'] ['approach'] ['local']
['bayesian'] ['probability'] ['mixture'] ['distribution'] ['parameters'] ['data'] ['models'] ['gaussian'] ['model'] ['likelihood']
['space'] ['linear'] ['vectors'] ['component'] ['matrix'] ['data'] ['analysis'] ['components'] ['information'] ['vector']
['random'] ['rate'] ['error'] ['examples'] ['gradient'] ['convergence'] ['algorithms'] ['time'] ['algorithm'] ['learning']
['language'] ['sequence'] ['nodes'] ['tree'] ['representations'] ['structure'] ['representation'] ['rules'] ['node'] ['rule']
['computer'] ['results'] ['level'] ['vector'] ['information'] ['number'] ['system'] ['memory'] ['parallel'] ['performance']
['neural'] ['vlsi'] ['voltage'] ['current'] ['output'] ['figure'] ['input'] ['chip'] ['analog'] ['circuit']
['cortical'] ['figure'] ['neurons'] ['cortex'] ['activity'] ['visual'] ['input'] ['cell'] ['model'] ['cells']
['tangent'] ['matching'] ['vectors'] ['feature'] ['set'] ['pattern'] ['character'] ['distance'] ['recognition'] ['image']
['theory'] ['generalization'] ['large'] ['field'] ['case'] ['learning'] ['energy'] ['function'] ['order'] ['noise']
```

The most similar titles to document 0 are,

```
['Observability of Neural Network Behavior ']
['Noisy Neural Networks and Generalizations,']
['A Precise Characterization of the Class of Languages Recognized by Neural Nets under Gaussian and Other Common Noise Di
']
['Analog Neural Networks of Limited Precision I: Computing with Multilinear Threshold Functions ']
['The Hopfield Model with Multi-Level Neurons ']
['On Properties of Networks of Neuron-Like Elements ']
['Complexity of Finite Precision Neural Network Classifier ']
['On the Effect of Analog Noise in Discrete-Time Analog Computations, ']
['Are Hopfield Networks Faster than Conventional Computers ?, ']
['On the Power of Neural Networks for Solving Hard Problems ']
```


Joint log likelihood code,

```
from scipy.special import loggamma
import numba_scipy
from numba import jit

@jit(nopython=True)
def joint_log_lik(doc_counts, topic_counts, alpha, gamma):
    """
    Calculate the joint log likelihood of the model

    Args:
        doc_counts: n_docs x n_topics array of counts per document of unique topics
        topic_counts: n_topics x alphabet_size array of counts per topic of unique words
        alpha: prior dirichlet parameter on document specific distributions over topics
        gamma: prior dirichlet parameter on topic specific distribuitons over words.

    Returns:
        ll: the joint log likelihood of the model
    """

    n_docs = doc_counts.shape[0]
    n_topics = doc_counts.shape[1]
    alphabet_size = topic_counts.shape[1]

    ll = 0

    for j in range(n_docs):
        term_1 = 0
        for i in range(n_topics):
            term_1 += loggamma(doc_counts[j][i] + alpha)
```

```
            term_2 = 0
            for i in range(n_topics):
                term_2 += doc_counts[j][i] + alpha
            term_2 = loggamma(term_2)
            ll += (term_1 - term_2)

    for i in range(n_topics):
        term_1 = 0
        for r in range(alphabet_size):
            term_1 += loggamma(topic_counts[i][r] + gamma)

        term_2 = 0
        for r in range(alphabet_size):
            term_2 += topic_counts[i][r] + gamma
        term_2 = loggamma(term_2)
        ll += (term_1 - term_2)

    return ll
```

Sampler code,

```
import numpy as np
import numba as nb

@nb.jit(nopython=True)
def sample_topic_assignment(topic_assignment,
                           topic_counts,
                           doc_counts,
                           topic_N,
                           doc_N,
                           alpha,
                           gamma,
                           words,
                           document_assignment):
    """
    Sample the topic assignment for each word in the corpus, one at a time.

    Args:
        topic_assignment: size n array of topic assignments
        topic_counts: n_topics x alphabet_size array of counts per topic of unique words
        doc_counts: n_docs x n_topics array of counts per document of unique topics

        topic_N: array of size n_topics count of total words assigned to each topic
        doc_N: array of size n_docs count of total words in each document, minus 1

        alpha: prior dirichlet parameter on document specific distributions over topics
        gamma: prior dirichlet parameter on topic specific distributions over words.

        words: size n array of words
        document_assignment: size n array of assignments of words to documents

    Returns:
        topic_assignment: updated topic_assignment array
        topic_counts: updated topic counts array
        doc_counts: updated doc_counts array
        topic_N: updated count of words assigned to each topic
    """
    n_words = len(words)
    n_topics = len(topic_N)
    alphabet_size = topic_counts.shape[1]

    alphabet_size = topic_counts.shape[1]

    topic_assignment_updated = topic_assignment
    topic_counts_updated = topic_counts
    doc_counts_updated = doc_counts
    topic_N_updated = topic_N

    for i in range(n_words):
        # get relevant indices for the current word
        doc_idx = document_assignment[i]
        top_idx = topic_assignment[i]
        word_idx = words[i]

        # remove the current word from counts
        topic_counts_updated[top_idx][word_idx] -= 1
        doc_counts_updated[doc_idx][top_idx] -= 1
        topic_N_updated[top_idx] -= 1

        a = doc_counts_updated[doc_idx] + alpha
        b = (topic_counts_updated[:, word_idx] + gamma) / (topic_N_updated + alphabet_size * gamma)
        probabilities = a * b
        probabilities /= np.sum(probabilities)

        # sample new topic assignment
        new_top_idx = np.argmax(np.random.multinomial(1, probabilities))

        # update counts
        topic_assignment_updated[i] = new_top_idx
        topic_counts_updated[new_top_idx][word_idx] += 1
        doc_counts_updated[doc_idx][new_top_idx] += 1
        topic_N_updated[new_top_idx] += 1

    return topic_assignment_updated, \
           topic_counts_updated, \
           doc_counts_updated, \
           topic_N_updated
```